

Prognosis of prostate adenocarcinoma metastasis using gene activation profiles as an unexplored source of information

Christopher C Thompson

November 10, 2016

1 Definition

1.1 Project Overview

The prostate is a glandular organ of the male reproductive system that helps to control urinary and reproductive functions. According to the charity, Prostate Cancer UK, one in eight British men will be diagnosed with prostate adenocarcinoma (henceforth, 'prostate cancer') in their lifetime [1]. Men over 50 years of age are often subjected to routine digital examinations, and/or a urine test (called the Prostate Secreted Antigen, or 'PSA' test) for signs of prostate cancer. However the False Positive Rate for these tests remain high [2] and, as such, the gold standard diagnosis is the Gleason test. In brief, a series of small needle sized biopsies are taken from the patient's prostate gland. Each biopsy is processed and scored by a pathologist for signs of abnormal cell type and structure. Gleason grades ranging from 2 to 5 are considered not malignant, whereas scores ranging from 6-10 are considered malignant and provide an estimation of cancer severity [3].

Contrary to some types of cancer, malignancies that remain local within the prostate are rarely lethal (survival rate of 99%). However, if a malignancy born of the prostate undergoes distant metastasis (the process of cancer cell migration to other sites in the body), the 5-year survival rate drops to 28% [4]. Because of this discrepancy, many men opt for radical prostatectomy (surgical removal of the entire prostate) despite not knowing whether a metastasis will occur or not. While resection does eliminate the chance of a future metastasis, it results in high morbidity (*e.g.* inability to control urination, loss of sexual function, etc).

Unfortunately, there are currently no prognostic tests for prostate cancer metastasis. The patient data that is typically available at the time of diagnosis is not rich enough to accurately predict the likelihood of prostate cancer metastasis [2]. A model that would be able to predict whether an untreated malignancy is likely to remain within the prostate or will metastasize to distant sites would be an invaluable tool in the decision between prostatectomy or surveillance. To generate such a model, it is clear that a more distinguishing data set is required.

One potential solution to this problem is an RNA-seq profile. In brief, RNA-seq is a technique that reads and counts RNA sequences in a biological specimen. What is RNA? When a gene is activated in a cell, the DNA sequence is read (or 'transcribed' or 'expressed') into an RNA molecule. RNA molecules are then read into protein molecules that function in all manner of operations within the cell. By reading and quantifying the RNA molecules that exist in a sample, one may determine which genes have been activated, and to what degree. A gene count profile (or 'RNA-seq' profile) is the estimation of activation for each of the full set of known genes in a biological specimen.

In lay terms, the full set of RNA molecules in a cell can be thought of as its blueprint. And if two sets of blueprints were incredibly similar, one would expect resultant buildings to be similar as well. In contrast, while a skyscraper and a lakehouse are both considered buildings, they would likely originate from very different sets of blueprints. In the same way, as metastatic cancer cells behave in drastically different ways than non-metastatic cells, one would assume that their RNA-seq profiles would be inherently different.

This difference should be detectable by RNA-seq analysis, though it is unlikely that any single gene could distinguish metastasis from a local malignancy. The ultimate goal of this project is to determine the probability of prostate cancer metastasis from an RNA-seq profile, generated from a prostate biopsy taken during the Gleason grading procedure.

1.2 Problem Statement

The primary questions that this project aims to answer are :

- Can the risk of prostate cancer metastasis state be predicted from a gene activation (RNA-seq) profile?

- If so, what genes (individually or in concert) are important for this distinction?

The goal of this project is to design a model that predicts the risk of prostate cancer metastasis using the gene activation profile derived from a patient’s prostate biopsy, taken at the initial Gleason grading diagnosis phase.

To achieve this goal, it is likely that a significant feature reduction exercise will be necessary, as each RNA-seq profile quantifies expression of 20501 human genes. After feature reduction, a model will be generated to quantify the risk of prostate cancer metastasis (probability from 0 to 1). Finally, a function or application will be engineered that receives an RNA-seq profile as an input and outputs a prediction for future metastasis state.

1.3 Metrics

An appropriate metric for the assessment of the probability of a binary class prediction is the Logarithmic Loss (Log Loss) score.

The equation for log loss is :

$$\text{logloss} = -\frac{1}{N} \sum_{i=1}^N (y_i * \log(p_i) + (1 - y_i) * \log(1 - p_i))$$

where p represents an observation’s predicted probability ($0 < p < 1$) and y represents the dummy variable of the actual binomial class {0,1}.

The log loss function provides a penalty score for each predicted observation in relation to the difference between the actual class {0,1} and predicted probability (0:1). Predictions that are both incorrect and confident are punished harshly. For instance, if a model were to return a certain outcome for binary classification (0 or 1), and that prediction was false (1 or 0), then infinity would be returned. Thus, in practice, the statistical programs will cap predictions away from absolute 0 or 1 prior to log loss assessment. On the other hand, an ultra-conservative model that predicted 0.5 for every observation (effectively not taking either stance in classification) would have a benchmark log loss score of approximately 0.693147.

2 Analysis

2.1 Data Exploration

‘The Cancer Genome Atlas’ (TCGA) is a research consortium set up to curate clinical data from thousands of patient participants, covering an array of cancer types. The data provided includes basic clinical information as well as DNA and RNA sequencing of cancer biopsies. These data sets are updated as new information becomes available. Thus each download represents a snapshot in an evolving, longitudinal study.

While detailed genomic and RNA sequence data is control-accessed, pre-processed gene count data is publicly available. Data can be downloaded via the consortium portal or acquired into data frame format using a package in the R language. An R script was written to access the data sets and write them locally in a python-readable format. The versions stored in the project repository were current at the time of the report date.

The clinical data set contains 22 features, of which several are irrelevant (*e.g.* all prostate cancer patients are ‘male’). Of the features, three were relevant and would be known at or very near the time of presentation: age, PSA test score, and Gleason score. One feature that would also be known but eliminated for ethical reasons is patient ‘race’. While a higher proportion of Black or Afro-Caribbean men are diagnosed with prostate cancer, the reasons for this are not fully understood [5] and significant evidence suggests that race/ethnicity should not be used in cases of genetic / gene activation analysis [6].

The outcome variable for this project is contained in the clinical data set, which is ‘pathologyNstage’. This label is composed of ‘n0’ or ‘n1’, representing **local** versus **metastatic** cancer, respectively. The current percentage of metastatic cases is approximately 16% (Figure ??), though this percentage is likely to increase as the age of the study increases (see Reflection section for discussion).

When grouped by Gleason score, it is evident that metastasis rates increased with cancer severity (Figure 2). This is intuitive, yet clearly not sufficient to determine whether a specific cancer, regardless of Gleason score, will metastasize or not. To illustrate, cancers that have been rated at a Gleason score of ‘9’ are still more likely to belong to the ‘n0’ class than the metastasis class.




Figure 1: Frequency of metastasis state ('pathologyNstage') in the TCGA Prostate adenocarcinoma cohort. Most cases in the TCGA cohort are labeled as 'non-metastatic', however this study is currently ongoing and some may be updated as metastases are diagnosed. The frequency counts are accurate from the date of this report.

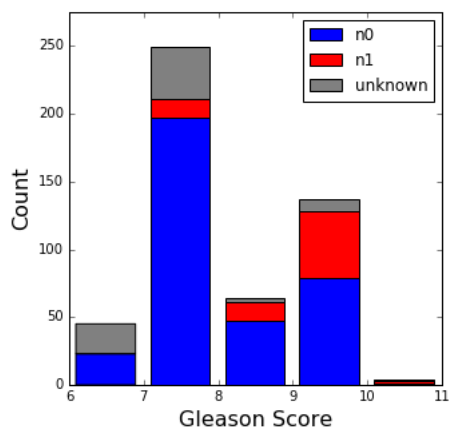


Figure 2: Frequency of metastasis state grouped by Gleason grade. The ratio of metastasis cases increases with cancer severity, as measured by Gleason grade.

2.2 Exploratory Visualization

The clinical information available at the onset of prostate cancer diagnosis is not rich enough to predict metastasis [2]. To corroborate this, age, PSA score, and Gleason grade were plotted in a scatter matrix in which each observation is colored based on metastasis state (Figure 3). While Gleason grade seems to correlate weakly to metastasis state, neither age or (surprisingly) PSA value were proportional to metastasis by visual analysis.

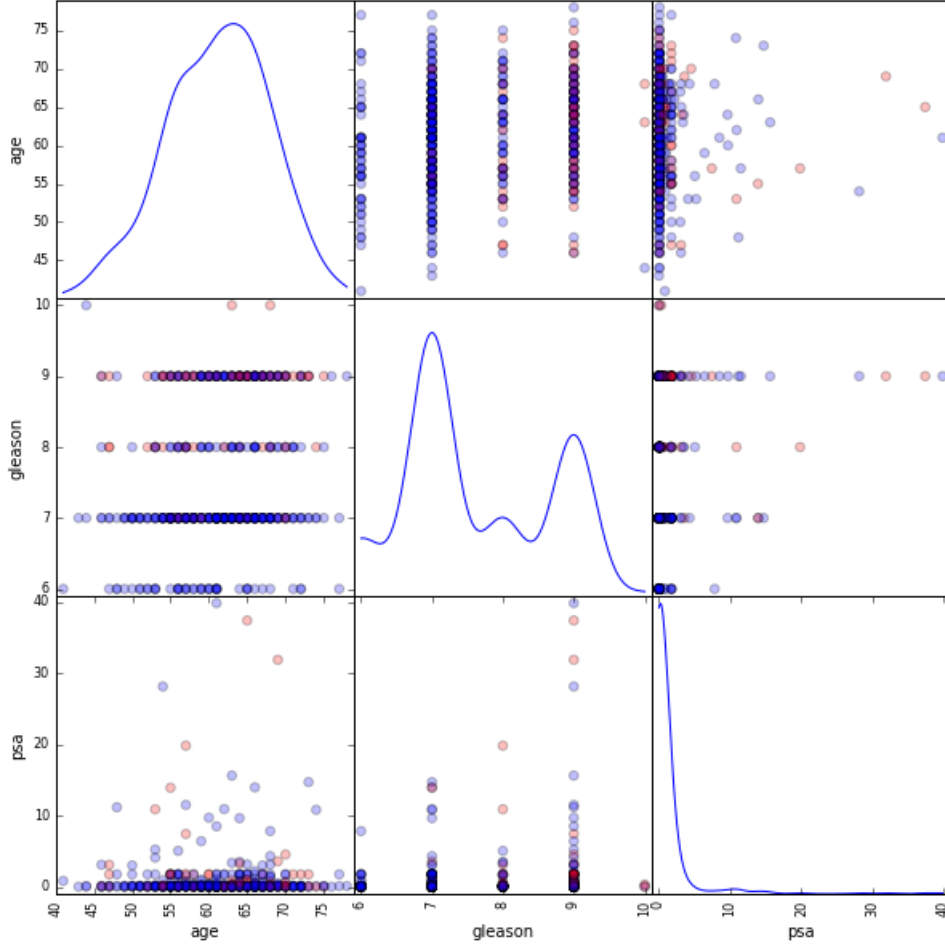


Figure 3: Relationship between age, PSA value, and Gleason grade in prostate cancer metastasis class ('n0' - blue, 'n1' - red). Only age appears to be distributed normally. Cases of metastasis appear to correlate weakly with increase in Gleason grade.

The primary data set to be used in this project is the gene count (or RNA-seq) matrix. This data set provides a value for gene expression level for every known human gene. The same patient index links the clinical data set to the gene count data set, of which 497 are common among the two. As a pilot experiment for the project rationale, an F-test was run for every gene feature in the normalized data set, comparing the 'n0' to 'n1' metastasis states. The results from this analysis are shown in Figure 4, and reveal that while most genes are not differentially expressed between metastasis states, some genes do appear to be (in-)activated in metastasis. This indicates that there are genes that could be used for predictive purposes and validates the project rationale.




Figure 4: Distribution of F-test statistics for the comparison of gene expression levels between the 'n0' and 'n1' metastasis states. Larger F-scores indicate a that the population means are different from one another, with respect to the variation observed within the whole cohort of samples. Clearly a moderate amount of genes exhibit differential activation levels in the 'n0' and 'n1' classes.

2.3 Algorithms and Techniques

The basic outline for project completion is as follows:

1. Feature selection (filter mechanism)
2. Feature compression into a lower dimensional set
3. Determine which projected features are important for metastasis state discrimination (wrapping mechanism)
4. Subset and train the probabilistic-classification algorithm
5. Measure performance of the trained algorithm on an independent validation ('test') set
6. Compare model performance to the benchmark model performance

The feature reduction exercise will utilize Random Forest Classifier, not as a classification algorithm, but as a method to measure the ability of sampled genes to separate the data set by metastasis class. By increasing the number of trees, the chance of sampling all features at least once is increased. Given noisy data, decision trees (and thus Random Forest) classifiers are prone to overfitting, so parameter limits on the tree depth and the minimum number of samples that can be split will be defined. The top portion of genes in 'Gini Importance' will be retained in a subset and carried into the next project phase.

The reduced feature data set will be compressed further using Principle Component Analysis (PCA). PCA is an unsupervised learning technique that transforms a data set into its principle components - *i.e.* the orthogonal vectors within the data that explain the greatest amount of its variance. By selecting the the most important components, several features may be combined into a lower number without significant loss of information. How many principle components will be carried into the algorithm training will depend on the amount of variance each component can explain. For example, if the first principle component that explains 95% of the data set variance, it would not be necessary to bring any other principle components forward for further analysis.

The probabilistic-classification algorithm chosen for this task is the logistic regression ('logit') model. This algorithm was chosen for its inherent ability to assess the probability of a binary outcome (*e.g.* metastasis or local malignancy) based on continuous input variables that cannot cleanly differentiate between class labels individually. Logistic regression classification is well suited for noisy data, in that it does not assume that there is any margin or hyperplane that is capable of separating class labels. Instead, it returns a likelihood of class assignment based on the linear combination of input variables as a single term into the logistic (or 'sigmoid') function:

$$P(class = 1|X) = \frac{1}{1 + e^{-X}}$$

where X is equated to :

Table 1: Benchmark logistic regression model coefficients for clinical features

Feature	Coefficient
age	-0.067344
PSA Value	0.025574
Gleason Grade	0.858936

$$X = \sum \beta_n * x_n$$

and x_n represents each feature, β_n represents each feature's coefficient, e is euler's number.

For algorithm training, a 'solver' is necessary to determine the optimal β coefficients to minimize penalties accumulated from a 'cost function'. There are several options for both the 'solver' and 'cost function', which also requires a regularization term, 'C'.

Initial Parameter Choices:

- Solver - The 'liblinear' solver is based on a coordinate descent algorithm and is ideal for small data sets.
- Cost function - the 'l2' cost function is more appropriate in situations where features have been pre-filtered or are few in number. In cases of high dimensionality, the 'l1' cost function should be preferred as it practically eliminates non-predictive features from contributing to the X term by minimizing the absolute values of the respective β coefficients.
- Regularization term - For situations of high noise (such as this), a larger term for C is recommended. However, the research plan was to optimize this term as needed with cross-validation. Thus it was left at the default value of 1 in the first training instance.

2.4 Benchmark

As personalized medicine (*e.g.* use of a patient's specific genetic or gene activation information for therapeutic decisions) has not been established in mainstream therapy, a benchmark for use of RNA-seq data for prognosis of metastasis was not available. Hypothetically, the most conservative model which predicts every test sample as having 50% chance of metastasis would yield a log loss score of 0.69314.

To establish a more fair benchmark for comparison, a logistic regression model (see methodology below) that incorporated the clinical information that would normally be known at the time of diagnosis was generated. These features were 'age', 'PSA score', and 'Gleason score'.




Figure 5: Visualization of a benchmark logistic regression predictive model performance. Training and test sets were stratified by Gleason grade, which appears to correlate to the benchmark model prediction of metastasis risk.

The coefficients for the three features in the model (representative values shown in Table 1) exhibited that Gleason score was by far the most predictive (approximately 0.85), and that age and, surprisingly, PSA score (which is the

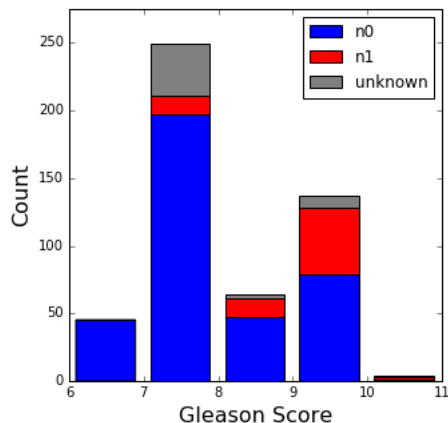


Figure 6: Distribution of known clinical features grouped by metastasis state (blue: 'n0', red: 'n1'). 'n0' was imputed for cases where label was missing and Gleason grade was mild.

current default test that doctors rely on for prostate cancer risk) provided very little use in classification. Figure 5 shows the relationship between Gleason grade and benchmark model probability of metastasis for each observation in the training (left) and test (right) sets. Data points are color-filled by actual metastasis state. The test set log loss score from this benchmark analysis was approximately 0.61 (using runs across 5 different seeds, 3), and thus could be considered marginally more useful than a '50% model'.

3 Methodology

3.1 Source Files

The data sets were retrieved from the TCGA portal using an R package, TCGA2STAT, and written to the local drive in feather format, which is python-readable. The R script used and feather files are available in this project's GitHub repository MLE_capstone. All algorithms were imported from the scikit-learn library, version 0.17.

3.2 Data Preprocessing

Samples with a Gleason score of 6 were homogenous in metastasis state (all 'n0'), though several cases were not labeled. In order to make more efficient use of the TCGA RNA-seq data set, 'n0' was imputed for all samples where no label existed and Gleason grade was defined as 6. The reasoning behind this decision was that those with low grade malignancy are usually not screened for metastasis and thus the lack of data label probably reflected the dispensibility of the metastasis test in such cases of mild malignancies. From a machine learning perspective, this imputation allowed more efficient use of a rather small data set. Given a much bigger data set, then this assumption would not be necessary, and all cases with missing label could be excluded. Indeed, for cases scored 7-10 on the Gleason scale where no labels were included in the clinical data were excluded from further analysis.

The gene count data retrieved from the TCGA portal was in an intermediary format. While the raw RNA-sequence reads had been processed into gene-level expression estimations, each profile required normalization for cross-sample comparison. Therefore, the initial gene count data set was transformed to transcripts per million (TPM) format. This dataframe was the base upon which further feature reduction and test train splitting would be performed.

3.3 Implementation

Feature reduction was completed in two steps. The first was to utilize the generation of a Random Forest Classifier to supply information regarding the importance of each gene in the separation of metastasis states. As the Random Forest model was not intended for actual classification purposes (as it was not optimal due to the small sample size of the data set), only key default parameters were altered. Specifically, the maximum tree depth was limited to 3 nodes, and the minimum number of samples that could be split was limited to 30. These parameter choices were intended to limit variance. The 'Gini Importance' of each feature was retrieved from the model and the genes ranked in the order of importance.

From this list, the original plan was to retain the top k -number of genes for PCA compression. However, run to run observation revealed that the set of genes was rarely identical. Many genes, such as *gne* were present in every case, however their ranking changed each run, which affected the subset composition in turn. To address the issue of feature stability, the Gini Importance selection process was repeated across 5 different random seeds, with genes kept only if present in the top k of each successive epoch. In this solution, k was set to 75 and resulted in a relatively stable selection of 12-15 genes. This subset was scaled to standard mean and unit variance using the sklearn Standard Scaler in preparation for further compression.

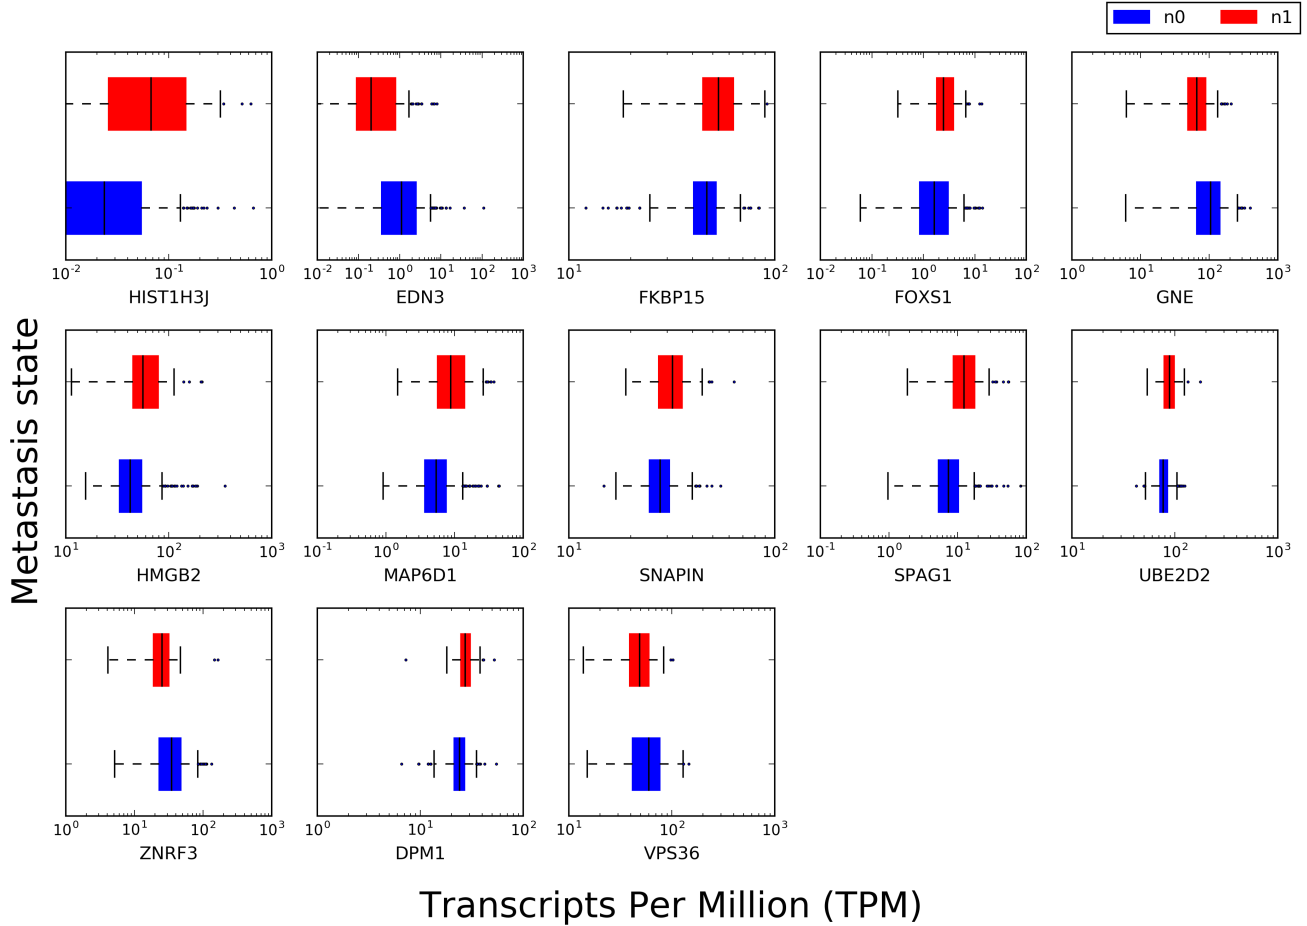


Figure 7: Genes with the highest 'Gini Importance' scores are generally not able of distinguishing metastasis class, individually.

The second phase of feature reduction was Principle Component Analysis compression. The PCA algorithm ranks a data set's orthogonal vectors by their variance, and rotates the data set to comply with the identified eigenvectors. The percent explained variance can be determined from associated eigenvalues in this process. Moreover, the contribution from each gene of the the k -gene subset can be determined and is shown for the first 3 PCs in Figure 8.

This 3-feature data set was then partitioned using the same indices from the first Train Test Split performed prior to the benchmark model generation. In detail, this split partitioned 70% of the samples into the training set, with 30% being held out for validation. The data was stratified by Gleason score, which was used as a surrogate measure for cancer severity. This decision was made to ensure that 'easy' (*e.g.* mild or extremely severe malignancies) and 'difficult' (*e.g.* malignancies on the border between moderate and severe) cases would be distributed equally. Another option would have been to stratify by metastasis label (see 'Reflection' section for label pair here discussion on this decision).

The training data set was then fed into a Logistic Regression Classifier model. For this learning, the class-weight parameter was set to 'balanced' in order to guard against confounding effects of the unbalanced label set in model performance. The regularization ('C') parameter was left at the default value of 1. The C term is inversely

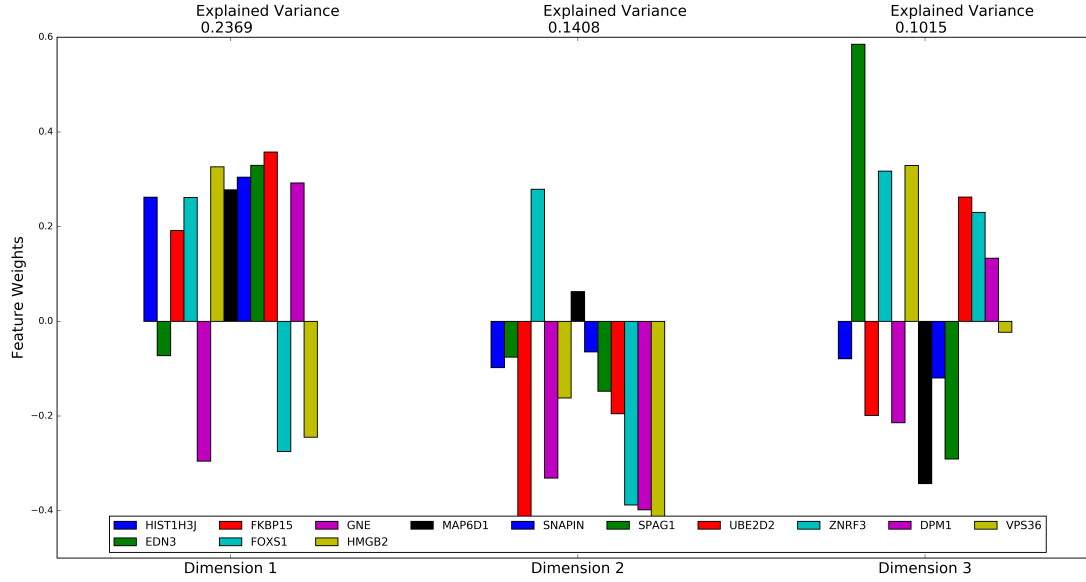


Figure 8: Explained variance and gene feature contribution to the first three principle components of the PCA transformation. Figure generated by code modified from the Udacity Machine Learning Nanodegree Project #3 - Creating Customer Segments.

proportional to the penalties awarded for misclassified samples. Hypothetically, a higher regularization term may have increased performance, however this was to be determined empirically in future optimizations.

Results were visualized using graphs generated with the matplotlib package. Performance of the logistic regression model was tested in all cases against the **held-out test set** using the log loss metric. For references, the $F\beta$ score ($\beta := 2$) and Matthews Correlation Coefficient scores are also listed, though they describe the performance of the algorithm to correctly classify metastasis state and do not measure performance in probabilistic prediction. Both the graphical analysis and metric reports were generated for each testing cycle using the scripts supplied in the 'Support Files' folder in the GitHub repository.

3.4 Refinement

In order to optimize the C parameter, a Logistic Regression CV classifier was generated using 4-fold cross-validation across a 10-log range for C. Performance was measured using log loss error rate. This process yielded an elevated term for C, 10000, indicating significant noise in the data set.

Because Gleason grade was clearly the most important clinical feature in predicting prostate cancer metastasis, it was added back to the training feature set to see if any improvement in model performance could be achieved.

4 Results

4.1 Model Evaluation and Validation

4.1.1 Final Model

The final logistic regression model receives 2 feature variables:

1. Gleason score
2. the first 3 PCs from a k -gene subset of RNA-seq expression values

The coefficients for Gleason grade and the first PC were routinely equivalent, indicating they contribute roughly evenly to dependent variable prediction. The 2nd and 3rd PCs did not contribute as much to the logistic regression

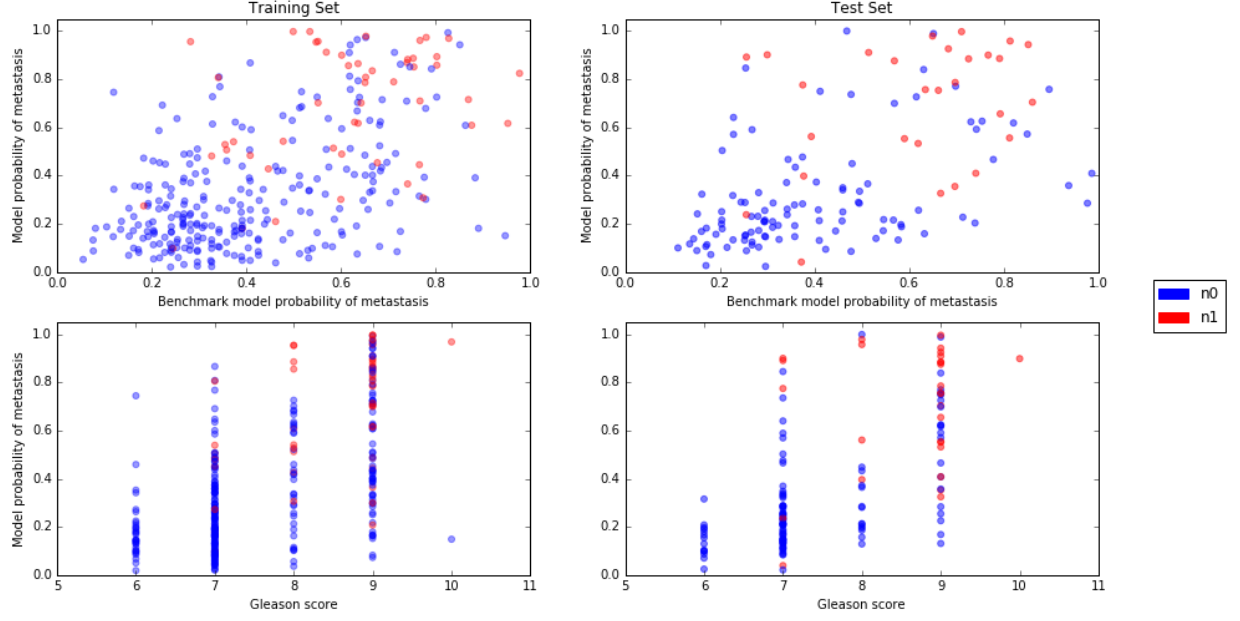


Figure 9: Model predictions on the Training and Test data sets after optimization of the regularization parameter.

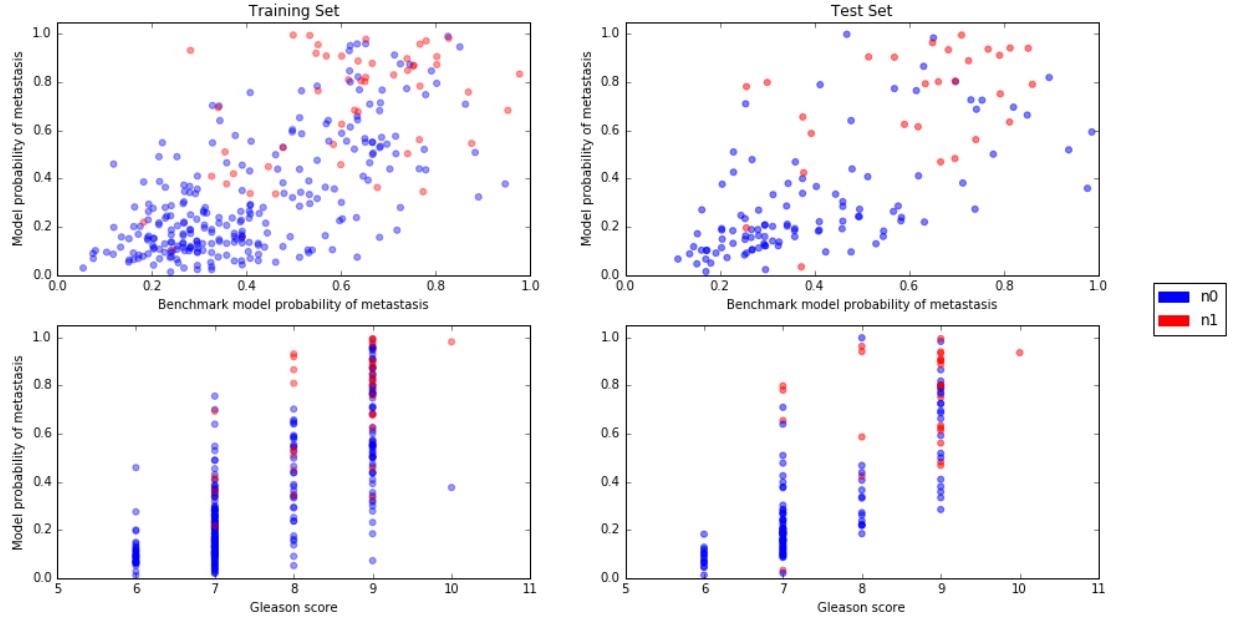


Figure 10: Model predictions after the addition of Gleason grade to the PC model.

decision function (see Table 2). The optimal regularization parameter was regularly determined as the maximum value tested, which is an indication of noisy (*i.e.* not linearly separable) data set.

Table 2: A representative logistic regression training outcome for the final model.

Feature	Coefficient
Gleason Grade	0.523239
First PC	0.603294
Second PC	-0.372170
Third PC	-0.190558

Table 3: Performance across 5 random seeds

Seed	Final Model LogLoss	Benchmark LogLoss	Improvement over Benchmark (%)
1	0.483871	0.613856	22.0
12	0.484437	0.617253	21.2
123	0.535816	0.627135	14.6
1234	0.484426	0.593156	18.3
12345	0.419407	0.574665	27.0

4.1.2 Test set Validation

This project’s strategy was to leave out 30% of the original data set to use as a true validation of the models’ generalization capability. The final model validation set log loss score ranged from 0.42 to 0.53 across five different random state seeds. Each value in this range was lower than the minimum benchmark score in the same 5 runs. Analyzed on a run by run basis (in which the training and test set cases are consistent), the final model achieved between 14.6 and 27.0% improvement over the benchmark.

4.2 Justification

The final logistic regression model performed better in predicting the probability of prostate cancer metastasis than the benchmark model in every run. Over the five consecutive runs described above, an average improvement of approximately 20% over the benchmark.

In order to test sensitivity of the model, a pipeline function was implemented that received RNA-seq profile and returned the final model probability of metastasis. To test the functionality of the pipeline application, all RNA-seq profiles where the label was missing and Gleason grade was 7-10 were subjected to prediction. Results from this analysis are shown in Figure 11. Clearly several of these cases were risk for metastasis according to the model.




Figure 11: Metastasis predictions for unlabeled TCGA cohort samples. TCGA cohort patient samples that did not include a metastasis label and were Gleason range 7-10 were omitted from model learning and validation. Samples are subjected to the risk analysis function and plotted against the benchmark model prediction (left) and Gleason score (right).

As a true test of sensitivity, matched patient benign controls were run through the pipeline function. These samples originated from areas of the prostate where no malignancy was evident (though malignancy was present within the same prostate gland in each case). For the gleason grade in these matched benign specimens, the grade

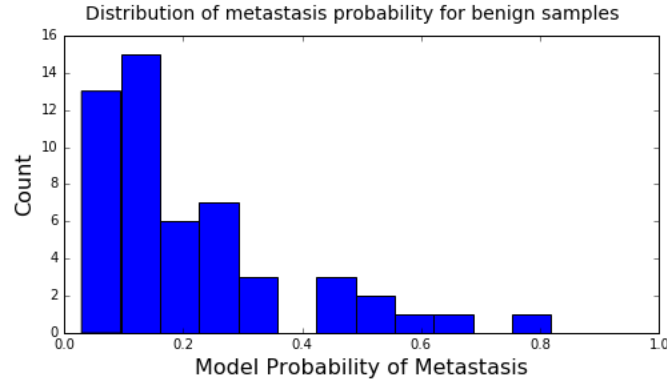


Figure 12: Analysis of risk from matched, benign controls from the TCGA cohort data reveal that the final model is stringent. Samples from this cohort were taken from benign areas of patient prostates where malignancies were present. Gleason grade was borrowed from the orthogonol malignancy. Despite this potential bias, the majority of samples are predicted with a low probability of metastasis.

from the malignant region was borrowed to make the test robust. For example, in some cases the RNA-seq of a benign specimen will be paired with a gleason score of 9, due to severe malignancy elsewhere in the prostate). An easier test would have been to impute a Gleason grade of 2-5 for each of these specimens. Nevertheless, despite the potential Gleason bias, the density of metastasis probability was right-skewed with the vast majority of predictions falling in the 0.05-0.3 range for the benign samples.

5 Conclusion

5.1 Free-Form Visualization

A summary of model improvement over the course of the run is shown in Figure 13.




Figure 13: Summary of the change in metric score over the optimization course of the project. The final model, that incorporates a 3 PCs from a k -gene subset of RNA-seq data, combined with Gleason grade performs better than the benchmark model in log loss error.

5.2 Reflection

5.2.1 Objective

The purpose of this project was to generate a model capable of supplying a patient and doctor with a metric for risk of prostate cancer metastasis that was more useful than the crude use of the 'Gleason Score'. To accomplish this, RNA-seq (gene activation profile) was explored as a potential inroad into personalized therapy for newly diagnosed prostate cancer patients. There were several issues that made this task difficult:

1. Small, wide sample data - the effective data set (containing Gene Activation profile and a metastasis label) was 446 samples by 20501 gene features.
2. 'Inaccurate' / 'Pre-mature' data labelling - The TCGA cohort is regularly updated and those listed as non-metastatic at the time of update could become metastatic at a later date. Indeed many of the 'non-metastatic' observations are still predicted to have a high chance of metastasis, despite many of the cases being used for training of the model algorithm (See Figure 15)
3. Noise in the data - no single gene or biomarker had been reported as capable of efficiently separating non-metastatic and metastatic cancers (corroborated in this project, Figure 7).

Thus from a machine learning perspective, it was clear from the project's onset that feature reduction and appropriate model selection would be paramount to success.

5.2.2 Feature Selection

There are many techniques for feature reduction. One avenue that was explored was feature elimination via a wrapping mechanism. However this approach was very slow and provided inconsistent results in which, and how many, features were important. A different approach was to utilize the training of an ensemble Random Forest classifier, not for its use in classification, but in order to access its assessment of which genes were most informative in separation of the metastasis classes. In order to stabilize the gene set to be retained, an iterative process was implemented in which the top 75 genes of a Random Forest to Gini Importance pipeline was compared to the list of the previous cycle. Those genes present in both lists were held over into a new set which was updated in each successive cycle. Thus any gene retained for the k -gene feature set would have been ranked in the top 75 of 20501 genes for Gini Importance in 6 separate Random Forest to Gini Importance cycles. This usually selected 12-15 genes for further compression in the next step. Interestingly, when visualized graphically, none of this reduced k -gene set could separate the metastasis state linearly (Figure 7), justifying selection of logistic regression as the eventual classification algorithm.

After PCA transformation of the k -gene subset, the plan was to provide the full complement of PCs (*i.e.* the number of dimensions from the k -gene set) to the logistic regression classifier as training data, and subsequently use each component's coefficient to assess which PCs were most able to explain the independent variable. However, graphical analysis of the principle component scatter matrix, grouped by metastasis state (Figure 14) curiously showed that the first principle component seemed to generate distinct gaussian distributions for each of the metastasis states, despite the fact that PCA is an unsupervised technique. This result was consistent independent of whether 5 through 500 genes were 'Gini' selected for PCA transformation.

How could this be? This result would be expected if a transformation technique such as linear discriminant analysis (LDA) had been employed, as LDA uses data label in order to determine the component vectors where class label is discriminated the most. PCA, on the other hand, is an unsupervised technique and had generated what appeared to be a discriminating component in the absence of label information. Upon reflection, it is perhaps not surprising that the eigenvector where the most variance in the data subset was contained (*i.e.* the first principle component) would separate the class labels, given that *only genes where a 'significant' difference in gene expression between the class labels* were supplied to the PCA fit and transform process.

By creating a pipeline from the Gini Importance filter directly into the PCA transformation, something similar to Linear Discriminant Analysis had been generated. Indeed, exploration of an supervised LDA compression of the k -feature set yielded a similar level of performance in the final model compared to compression via Gini Importance to PCA pipeline (Data not shown in the project's GitHub repository master branch, though this feature can be observed in terminal commit in the LDA_exploration branch).

The 3-component feature set taken from this transformation was split on the same indices that were generated in the training and validation sets used in the benchmark analysis. This was done to aid in model to model comparisons within each run.




Figure 14: Analysis of PCA transformation of a k -gene feature subset. The first principle component of PCA transformation separates metastasis state more efficiently than any single gene from the input set. The second and third principle components are also shown for reference.

5.2.3 Model Selection

Having completed a feature selection and compression technique, in which at least the first principle component seemed capable of distinguishing among metastasis class (graphical analysis, Figure 14), a logistic regression classifier was chosen as the predictive model. Logistic regression was preferred to other hyperplane-based techniques, such as support vector machines (SVM), due to the noise that was expected in the data set. SVM classifiers attempt to define the hyperplane by which the margin between the class labels is maximized. In situations where data is not easily separable, this result can be unstable, and at times, arbitrary. Moreover, SVM does not provide a true probability of class assignment, as was the objective of the project. In contrast, logistic regression assumes that no feature is capable of explaining the outcome variable completely, but that the combination of features should be able to provide an odds-probability of class assignment. These assumptions are consistent with the RNA-seq data set employed in this project. Moreover, as the objective of this project was to provide a probability of metastasis, the output of logistic regression classifier was perfectly suited.

5.2.4 Training and Optimization

Separate Train and Test indices were stratified based on cancer severity prior to the benchmark analysis and the final PCA compressed (3-components) were subset into these indices. A logistic regression classifier was trained and optimized on the Training set, prior to validation on the Test set. Additionally, Gleason grade was added as a feature to the model to determine whether an increased predictive performance could be achieved. Indeed, the addition of this feature provided an incremental decrease in log loss error.

The final release version of the code was run across 5 seeds and performance in the primary metric (log loss)




Figure 15: Metastasis prediction of samples labeled as non-metastatic. Despite being labeled as 'non-metastatic', some observations are predicted with high risk of metastasis.

(Table 3), and secondary metrics (F2 and MCC) were recorded, compared to the benchmark.

5.2.5 Model Performance

The mean improvement of the final model over the benchmark model per run was 20%. Considering the small starting sample size, unlabeled and mislabeled data, and number of features, this improvement should be considered a success. A high level of sensitivity was observed in the model, as evidenced by the assessment of benign specimens (Figure 12).

In every run tested, the performance of the final model exceeded performance of the benchmark model by at least 14% in log loss score. The pipeline exhibited in this project could be re-appropriated for other types of RNA-seq based classifications. By looking for individual genes whose activation level explain a certain condition, researchers may be missing the opportunity to provide valuable disease prognosis. Instead, by performing a feature selection and compression, researchers may be able to predict disease more regularly at the sacrifice of knowing *exactly* what genes are causal.

As mentioned, one issue with this data set is that the TCGA cohort study is continuously updated. The question of '*At what duration can a non-metastatic participant be confidently labeled as such?*'. Change in metastasis state can only move in one direction, thus hypothetically the performance of the final model should trend towards more accurate, as those non-metastatic cases that were predicted to metastasize are updated in the study information. Figure 15 shows that a significant portion of 'n0' cases belong in this category.

Unfortunately in the context of the TCGA cohort study, the link between patient and barcode has been broken for ethical reasons, meaning that such patients with high risk can not be identified for extra care and surveillance.

5.3 Improvement

While certainly prosaic, this project could use an increased number of participants. With the short and wide input data set, the options for algorithm choice were severely limited. Convergence between the Training and test set error rates in log loss score indicate variance is not an issue, though the low number of final input features could very well be introducing bias.

5.4 Final Remarks

The objective of this machine learning exercise was to develop a model that incorporated RNA-seq profiles (a feature set of 20501 genes) into the prediction of prostate cancer metastasis risk. Given the small sample size and incomplete / inaccurate labelling of noisy data, a 15-20% improvement over the benchmark rate should be considered a success, especially considering that many 'non-metastatic' cases may eventually be re-labeled as metastatic. Increased study could be warranted to determine if use of this model could benefit those diagnosed with prostate cancer in the very difficult decision of resection versus surveillance. Finally, it may be possible to utilize this same pipeline for metastasis predictions in other types of cancer simply by modifying the input data set.

6 References

References

- [1] Prostate Cancer UK, <http://prostatecanceruk.org/prostate-information>, Accessed 01-August-2016.
- [2] Brawley OW, Thompson IM Jr, Grnberg H. (2016) *Evolving Recommendations on Prostate Cancer Screening*, **Am Soc Clin Oncol Educ Book**, 35, e80-7.
- [3] Humphrey PA. (2004) *Gleason grading and prognostic factors in carcinoma of the prostate*, **Mod Pathol**. 17, pp 292-306.
- [4] Cancer.org, <http://www.cancer.org/cancer/prostatecancer/detailedguide/prostate-cancer-survival-rates>, accessed 01-August-2016.
- [5] Shea PR, Ishwad CS, Bunker CH, Patrick AL, Kuller LH, Ferrell RE. (2008) *RNASEL and RNASEL-inhibitor variation and prostate cancer risk in Afro-Caribbeans.*, **Prostate**, 68, pp 354-9.
- [6] Yudell M, Roberts D, Desalle R, & Tishkoff S. (2016) *Taking race out of human genetics*, **Science** 351, pp 564-565.